

## » Convolutional Networks

- \* *ConvNets, CNNs* → have driven change in image processing since around 2012
- \* Usually *Deep learning* = *ConvNets*
- \* Also increasingly being used for text analytics/natural language processing, and other places where input data is a sequence.
- \* ConvNets consist of input and output layers plus multiple hidden layers, e.g. 10-100 layers. Hence “deep” learning since MLPs (“shallow” networks) usually have just one hidden layer.
- \* Each layer takes output of previous layer as its input.
  - \* Main types of layer: *Convolutional*, *pooling*, *fully connected*
- \* Some good resources online (also plenty of terrible ones) e.g.
  - \* <https://www.coursera.org/learn/convolutional-neural-networks/>
  - \* Stanford CS231 course <https://cs231n.github.io/>

## » Convolutional Layer

- \* Nodes in a convolutional layer use a *kernel* or *filter*
- \* Basic primitive: take a matrix as input, apply kernel to it (*convolve* the matrix and kernel) and produce a matrix as output
- \* Conventional to use \* to denote convolution, try not to mix up with multiplication
- \* Example:

1	2	3	4	5
1	3	2	3	10
3	2	1	4	5
6	1	1	2	2
3	2	1	5	4

Input

\*

1	0	-1
1	0	-1
1	0	-1

Kernel

=


Output

## » Convolutional Layer

1 <sup>1</sup>	2 <sup>0</sup>	3 <sup>-1</sup>	4	5
1 <sup>1</sup>	3 <sup>0</sup>	2 <sup>-1</sup>	3	10
3 <sup>1</sup>	2 <sup>0</sup>	1 <sup>-1</sup>	4	5
6	1	1	2	2
3	2	1	5	4

Input

1	0	-1
1	0	-1
1	0	-1

Kernel

-1		

Output

$$* \quad 1 \times 1 + 1 \times 1 + 3 \times 1 + 2 \times 0 + 3 \times 0 + 2 \times 0 + 3 \times (-1) + 2 \times (-1) + 1 \times (-1) = 5 - 6 = -1$$

## » Convolutional Layer

1	2 <sup>1</sup>	3 <sup>0</sup>	4 <sup>-1</sup>	5
1	3 <sup>1</sup>	2 <sup>0</sup>	3 <sup>-1</sup>	10
3	2 <sup>1</sup>	1 <sup>0</sup>	4 <sup>-1</sup>	5
6	1	1	2	2
3	2	1	5	4

Input

\*

1	0	-1
1	0	-1
1	0	-1

Kernel

=

-1	-4	

Output

\*  $2 \times 1 + 3 \times 1 + 2 \times 1 + 3 \times 0 + 2 \times 0 + 1 \times 0 + 4 \times (-1) + 3 \times (-1) + 4 \times (-1) = 7 - 11 = -4$

## » Convolutional Layer

1	2	3 <sup>1</sup>	4 <sup>0</sup>	5 <sup>-1</sup>	*
1	3	2 <sup>1</sup>	3 <sup>0</sup>	10 <sup>-1</sup>	
3	2	1 <sup>1</sup>	4 <sup>0</sup>	5 <sup>-1</sup>	
6	1	1	2	2	
3	2	1	5	4	

Input

1	0	-1
1	0	-1
1	0	-1

Kernel

=

-1	-4	-14

Output

\*  $3 \times 1 + 2 \times 1 + 1 \times 1 + 4 \times 0 + 3 \times 0 + 4 \times 0 + 5 \times (-1) + 10 \times (-1) + 5 \times (-1) = 6 - 25 = -14$

## » Convolutional Layer

1	2	3	4	5
<sup>1</sup> 1	<sup>0</sup> 3	<sup>-1</sup> 2	3	10
<sup>1</sup> 3	<sup>0</sup> 2	<sup>-1</sup> 1	4	5
<sup>1</sup> 6	<sup>0</sup> 1	<sup>-1</sup> 1	2	2
3	2	1	5	4

Input

$$\begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline 1 & 0 & -1 \\ \hline \end{array}$$

Kernel

$$\begin{array}{|c|c|c|} \hline -1 & -4 & -14 \\ \hline 6 & & \\ \hline & & \\ \hline \end{array}$$

Output

- \*  $1 \times 1 + 3 \times 1 + 6 \times 1 + 3 \times 0 + 2 \times 0 + 1 \times 0 + 2 \times (-1) + 1 \times (-1) + 1 \times (-1) = 10 - 4 = 6$
- \* Observe that applying  $3 \times 3$  kernel to a  $5 \times 5$  matrix gives a  $3 \times 3$  matrix

## » Convolutional Layer

1	2	3	4	5
1	3	2	3	10
3	2	1	4	5
6	1	1	2	2
3	2	1	5	4

Input

\*

1	0	-1
1	0	-1
1	0	-1

Kernel

=

-1	-4	-14
6	-3	-13
9	-6	-8

Output

## » Example: Edge Detection



\* Suppose we want to detect vertical edges in image ...

\* Try kernel:

1	0	-1
1	0	-1
1	0	-1



## » Example: Edge Detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

 $*$ 

1	0	-1
1	0	-1
1	0	-1

 $=$ 


## » Example: Edge Detection

$10^{-1}$	$10^0$	$10^{-1}$	0	0	0
$10^{-1}$	$10^0$	$10^{-1}$	0	0	0
$10^{-1}$	$10^0$	$10^{-1}$	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

 $\ast$ 

1	0	-1
1	0	-1
1	0	-1

 $=$ 

0			

\*  $10 \times 1 + 10 \times 1 + 10 \times 1 + 10 \times 0 + 10 \times 0 + 10 \times 0 + 10 \times (-1) + 10 \times (-1) + 10 \times (-1) = 0$

## » Example: Edge Detection

10	$10^1$	$10^0$	$0^{-1}$	0	0
10	$10^1$	$10^0$	$0^{-1}$	0	0
10	$10^1$	$10^0$	$0^{-1}$	0	0
10	10	10	0	0	0
10	10	10	0	0	0

 $\ast$ 

1	0	-1
1	0	-1
1	0	-1

 $=$ 

0	30		

$$\ast \quad 10 \times 1 + 10 \times 1 + 10 \times 1 + 10 \times 0 + 10 \times 0 + 10 \times 0 + 10 \times (-1) + 0 \times (-1) + 0 \times (-1) = 30$$

## » Example: Edge Detection

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

 $*$ 

1	0	-1
1	0	-1
1	0	-1

 $=$ 

0	30	30	0
0	30	30	0
0	30	30	0

- \* Observe that non-zero values in output highlight the edge

## » Example: Edge Detection

Dark→light vs light→dark edges

0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

 \* 

1	0	-1
1	0	-1
1	0	-1

 = 

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0

- \* Sign of output depends on whether transition is dark→light or light→dark

- \* To detect horizontal edges use kernel:

1	1	1
0	0	0
-1	-1	-1

- \* Similarly 45° angled edges, 70° etc
- \* Can use kernel larger than  $3 \times 3$ , but  $3 \times 3$  is v common in ConvNets.

## » Learning Kernels

- \* Rather than hand-crafting kernels, *learn* them !

0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

 $*$ 

$w_1$	$w_2$	$w_3$
$w_4$	$w_5$	$w_6$
$w_7$	$w_8$	$w_9$

 $=$ 

$x_1$	$x_2$	$x_3$	$x_4$
$x_5$	$x_6$	$x_7$	$x_8$
$x_9$	$x_{10}$	$x_{11}$	$x_{12}$

- \* Output from convolution node is a matrix, we need to map this to a scalar output/prediction  $\rightarrow$  reshape/flatten matrix into a vector  $x = [x_1, x_2, \dots, x_{12}]$ , then use this as input to one of our usual models E.g.
  - \* Linear model  $\hat{y} = \theta^T x$
  - \* NB: Can add more convolution & other layers before map to output  $\rightarrow$  will come back go this soon!
- \* Use training data to learn the unknown (i) output parameters  $\theta$  and (ii) kernel weights  $w = [w_1, w_2, \dots, w_9]$ :
  1. Define cost function
  2. Use gradient descent  $\rightarrow$  typically use stochastic gradient descent variant.
- \* Back in familiar territory: a model mapping from input (matrix of pixel values) to predicted output, model has unknown parameters, learn these using a cost function+training data.

## » Padding

- \* Applying  $3 \times 3$  kernel to  $5 \times 5$  matrix gives  $3 \times 3$  matrix  $\rightarrow$  output is *smaller* than input
- \* Often want to keep output same size as input  $\rightarrow$  use padding

0	0	0	0	0	0	0
0	1	2	3	4	5	0
0	1	3	2	3	10	0
0	3	2	1	4	5	0
0	6	1	1	2	2	0
0	3	2	1	5	4	0
0	0	0	0	0	0	0

 $\times$ 

1	0	-1
1	0	-1
1	0	-1

Kernel

 $=$ 


- \* Add extra rows and columns of zeros to pad original  $5 \times 5$  input matrix out to  $7 \times 7$ . Apply kernel to this padded matrix to obtain  $5 \times 5$  output i.e. output same size as original  $5 \times 5$  input.
- \* Some terminology:
  - \* *Valid* convolution: apply kernel directly to input  $\rightarrow$  output is smaller than input
  - \* *Same* convolution: pad original input then apply kernel  $\rightarrow$  output is same size as input

## » Strided Convolutions

- \* In previous examples we moved kernel along by one column/row at each step → we used a stride of 1
- \* Can also use larger strides e.g. stride 2

1	2	3	4	5
1	3	2	3	10
3	2	1	4	5
6	1	1	2	2
3	2	1	5	4

 \* 

1	0	-1
1	0	-1
1	0	-1

 = 

-1	

1	2	3	4	5
1	3	2	3	10
3	2	1	4	5
6	1	1	2	2
3	2	1	5	4

 \* 

1	0	-1
1	0	-1
1	0	-1

 = 

-1	-14

1	2	3	4	5
1	3	2	3	10
3	2	1	4	5
6	1	1	2	2
3	2	1	5	4

 \* 

1	0	-1
1	0	-1
1	0	-1

 = 


-1	-14
9	

- \* Observe that increasing stride reduces the size of the output matrix




## » Multiple Channels

- \* Gray-scale images are described by a single matrix, each element of matrix specifying shade of corresponding pixel
- \* Colour images are described by *three* matrices.
  - \* E.g. RGB image has one matrix specifying red intensity of each pixel, one specifying green and one specifying blue.





## » Multiple Channels



- \*  $32 \times 32 \times 3$  input has three *channels*, each channel is a  $32 \times 32$  matrix of values
  - \* The  $32 \times 32 \times 3$  stack of three  $32 \times 32$  matrices is called a *tensor*
- \* We define a separate  $3 \times 3$  kernel for each channel, so overall have a  $3 \times 3 \times 3$  kernel
- \* How to calculate output?

## » Multiple Channels



- \* Apply each  $3 \times 3$  kernel to its corresponding  $32 \times 32$  input channel. This gives three  $30 \times 30$  output matrices.
- \* Now add element (1,1) of each of the three matrices together to get element (1,1) of final output. Repeat for all elements  $(i,j)$ ,  $i = 1, \dots, 30, j = 1, \dots, 30 \rightarrow$  end result is a single  $30 \times 30$  matrix as output
- \* Note: number of channels of kernel *must* match number of input channels. E.g. if have 3 input channels then need 3 kernel channels.

## » Multiple Channels

More detailed example:

- \* Three  $3 \times 3$  kernels:  
channel 1

1	0	-1
1	0	-1
1	0	-1

channel 2

2	0	-2
2	0	-2
2	0	-2

channel 3

3	0	-3
3	0	-3
3	0	-3

- \* Three input channels:  
channel 1

1	2	3	4
1	3	2	3
3	2	1	4
6	1	1	2

channel 2

4	3	1	2
2	6	1	2
1	3	1	2
3	2	1	0

channel 3

7	2	3	4
6	3	2	3
5	2	1	4
1	5	6	4


- \* Output is obtained by applying channel  $i$  kernel to channel  $i$  input then summing. E.g. element (1,1) of output is

$$\begin{aligned} & 1 \times 1 + 1 \times 1 + 3 \times 1 + 2 \times 0 + 3 \times 0 + 2 \times 0 + 3 \times (-1) + 2 \times (-1) + 1 \times (-1) \\ & + 4 \times 2 + 2 \times 2 + 1 \times 2 + 3 \times 0 + 6 \times 0 + 3 \times 0 + 1 \times (-2) + 1 \times (-2) + 1 \times (-2) \\ & + 7 \times 3 + 6 \times 3 + 5 \times 3 + 2 \times 0 + 3 \times 0 + 2 \times 0 + 3 \times (-3) + 2 \times (-3) + 1 \times (-3) \\ & = -1 + 8 + 36 = 43 \end{aligned}$$

- \* Now shift all kernels by one column and repeat to get element (1,2) of output, and so on.

## » Multiple Filters

- \* We can apply several filters to the same input and stack their outputs together
- \* E.g. Apply two  $3 \times 3 \times 3$  kernels to a  $32 \times 32 \times 3$  input to get a  $30 \times 30 \times 2$  output:



- \* Note: number of output channels can be larger/smaller/same as number of input channels
- \* Kernel weights  $w$ , input  $a$ . After convolution output is  $w * a$ . Here  $a$  and  $a * w$  are both tensors i.e. a stack of matrices.
- \* What is number of parameters  $w$  in this setup?
  - \* Each  $3 \times 3 \times 3$  kernel has 27 weights
  - \* One  $3 \times 3 \times 3$  kernel for each output channel, so  $27 \times 2 = 54$  weights.